

When AI turns grant evaluation into a lottery

Martin Bulla^{1,✉} and Peter Mikula²

(bullab, Faculty of Environmental Sciences, Czech University of Life Sciences Prague, Czechia)

¹ bullam@fzp.czu.cz; ORCID: 0000-0002-1328-1927

² mikulap@fzp.czu.cz; ORCID: 0000-0002-2731-9105

Research funding schemes are increasingly becoming zero-sum games, characterised by a rising volume of applications chasing stagnant and often decreasing pools of capital^{1,2}. Despite the immense resources diverted toward peer review, traditional scoring systems are struggling to reliably distinguish scientific merit³⁻⁵. Here, we argue that the widespread adoption of generative AI (specifically Large Language Models) over the past year has accelerated this decline, pushing the current funding model toward 'excellence saturation', a state in which conventional scoring systems further lose their discriminative power. This raises fundamental questions about how funding decisions are ultimately made.

To illustrate, the latest EU *Marie Skłodowska-Curie Actions* postdoctoral fellowship call saw a nearly twofold increase in applications and striking score compression, with only ~5% of proposals fell below the 70% threshold, compared with ~20% in previous years⁶ (Fig. 1). This shift is unlikely to reflect a sudden population-wide increase in researcher excellence; rather, it indicates the proliferation of ever more capable AI tools that substantially improve structure, clarity, and persuasiveness of grant proposals, effectively decoupling a proposal's form from the underlying quality and innovativeness of research ideas^{7,8}. Moreover, beyond stylistic refinement, AI tools can explicitly optimise proposals against published evaluation criteria, generating text that mirrors scoring rubrics and systematically addresses reviewer expectations. This creates a 'Red Queen' race in grant writing⁷. If every applicant leverages AI to polish their proposal to a 95% standard, the baseline for perceived excellence inevitably shifts upward. Reviewers are then forced into a cycle of 'hair-splitting', attempting to differentiate between 'perfect' and 'more perfect' applications. In this environment, the signal-to-noise ratio collapses, and minute score differences no longer track meaningful distinctions in scientific merit^{3-5,9}. Indeed, when ~20% of proposals score above 95% (Fig. 1b)⁶, funding decisions resemble a lottery.

Furthermore, the current situation threatens to exacerbate existing inequalities in science. While generative tools were initially hailed as a way to level the playing field for non-native English speakers¹⁰, they may instead install a 'Silicon Ceiling'^{8,10}. Wealthy teams can afford premium, fine-tuned AI suites or professional editors to 'humanize' and optimize AI drafts, while researchers in resource-constrained regions must rely on limited, free versions¹⁰. When funding success depends on a level of polish that is now commercially available, the system risks rewarding those with the best tools rather than those with the best ideas^{7,9}.

This reflects a structural breakdown of fine-grained peer review under conditions of score saturation. The opportunity costs for applicants and reviewers in the hyper-competitive schemes now often exceed the benefits of marginal improvements in selection accuracy, effectively draining resources from research itself^{1,2}. If AI-driven saturation renders traditional scoring obsolete, funding bodies need to shift from fine-grained ranking of excellence toward managing an abundance of excellence. Adopting alternative allocation mechanisms, such as lotteries for high-scoring proposals that pass a rigorous quality threshold, is no longer a radical suggestion but a necessary response to the automated era of grant writing³⁻⁵.

Competing interests

The authors declare no competing interests.

Acknowledgements.

We were supported by the Research Excellence in Environmental Sciences grant from the Faculty of Environmental Sciences, Czech University of Life Sciences Prague (REES 003 to MB).

References

1. Dresler, M. *et al. Nat Hum Behav* **6**, 607–608 (2022).
2. Schweiger, G. *Nature* <https://doi.org/10.1038/d41586-025-04060-x> (2025) doi:10.1038/d41586-025-04060-x.
3. Fang, F. C., Bowen, A. & Casadevall, A. *eLife* **5**, e13323 (2016).
4. Schweiger, G. *et al. Proceedings of the National Academy of Sciences* **121**, e2407644121 (2024).
5. Gigerenzer, G. *et al. Proceedings of the National Academy of Sciences* **122**, e2401237121 (2025).
6. European Commission. *EU Funding & Tenders Portal* <https://ec.europa.eu/info/funding-tenders/opportunities/portal/screen/opportunities/topic-details/HORIZON-MSCA-2025-PF-01-01> (2026).
7. Doshi, A. R. & Hauser, O. P. *Science Advances* **10**, eadn5290 (2024).
8. Hao, Q., Xu, F., Li, Y. & Evans, J. *Nature* **649**, 1237–1243 (2026).
9. Roumbanis, L. *Science, Technology, & Human Values* **44**, 994–1019 (2019).
10. Amano, T. & Berdejo-Espinola, V. *Trends in Ecology & Evolution* **40**, 273–285 (2025).




Fig. 1 | Temporal shift and score compression in evaluations of Marie Skłodowska-Curie Actions postdoctoral fellowships. **a**, Temporal trend in the percentage of proposals scoring at or above threshold across evaluation panels. The 2025 cohort (yellow) shows a marked divergence from the tightly clustered historical distribution (2018–2024), shifting toward higher scores across all deciles. For example, the 85% ‘Seal of Excellence’ threshold was reached by ~70% of proposals in 2025, compared to ~40% in previous years (medians). **b**, Temporal trend of “excellence saturation”, defined as the percentage of proposals achieving a score $\geq 95\%$. **a-b**, The sharp divergence in 2024–2025 is temporally aligned with the rapid institutional adoption of Large Language Models in research workflows⁸. Data extracted from the [EU Funding & Tenders Portal](https://ec.europa.eu/info/funding-tenders/opportunities/portal/screen/opportunities/topic-details/HORIZON-MSCA-2025-PF-01-01) and available at https://github.com/MartinBulla/correspondance_funding.