

🔒 Post blocked



Thomas Dietterich
@tdietterich.bsky.social

Follow

We are implementing a similar policy at [@arxiv.bsky.social](#). If there is incontrovertible evidence of LLM slop in a paper, this means the authors did not take the time to read the LLM output and we can't trust anything else in the paper. Penalty is 1 year ban from arXiv followed by...

12:26 PM · May 13, 2026

139 reposts 35 quotes 412 likes 24 saves

💬 15

↻ 174

❤️ 412



Social Indicators Research @socialindicators.bsky.social · 4mo

We are doing it as well; permanent ban.

12 authors banned so far this year.

Not a single one has replied asking to challenge it, to appeal. I don't think we're being overzealous.

💬 1

↻ 4

❤️ 31



Social Indicators Research @socialindicators.bsky.social · 4mo

Think of it as aligned to the interests of "genuine" authors, the people doing their work the right way.

Peer review is a scarce resource -- it should never be expended on fraudulent work.

💬

↻ 2

❤️ 16



Thomas Dietterich @tdietterich.bsky.social · 4mo

a requirement that subsequent arXiv submissions must first be accepted at a reputable peer-reviewed venue.

💬 3

↻ 5

❤️ 80



Alexander Helber @ahelber.bsky.social · 4mo

Will there also be naming and shaming? Seems like introducing some reputation



Bluesky

Create account

Sign in



Ekaterina (Katja) Pashevich @cosmickate.bsky.social · 4mo

I second the shaming 😊



Peter Coles 🇷🇺 🇧🇪 🇮🇪 @telescoper.bsky.social · 4mo

There doesn't appear to be anything about this in the actual Code of Conduct on the arXiv website

info.arxiv.org/help/policie...

Code of conduct - arXiv info | arXiv e-print repository

🌐 info.arxiv.org



Rod Van Meter @rdvquantum.bsky.social · 4mo

Yeah, I looked, too, and don't see anything on the website yet.



David Picard @davidpicard.eurosky.social · 4mo

The classifier that says if there is incontrovertible evidence of LLM slop in a paper, is it open source? I'd like to see how it's made and also I'd like to use it.




Jon Ullman @thejonullman.bsky.social · 4mo

Is LLM slop defined precisely enough for there to be incontrovertible evidence of it?




Iris van Rooij 🌐 @irisvanrooij.bsky.social · 4mo

irisvanrooijcogsci.com/2025/08/12/a...



Bluesky

Sign in



AI slop and the destruction of knowledge

Cite as: van Rooij, I. (2025) AI slop and the destruction of knowledge. This week I was looking for info on what cognitive scientists mean when they speak of 'domain-general' cognition. I was ...

irisvanrooijcogsci.com

1

2

12

Bookmark Share ...



Clément Canonne @ccanonne.github.io · 4mo

I reckon Jon is asking specifically "what is the criterion to classify sthg as 'AI slop'", esp. at the scale of arXiv. (I don't see such a definition in the linked article?)

(Reading [@tdietterich.bsky.social](#)'s post, it's unclear if the policy applies to only hallucinated refs or "AI slop" broadly)

2

2

9

Bookmark Share ...



Iris van Rooij @irisvanrooij.bsky.social · 4mo

It will require expert cognition to detect slop, that was my point.

2

2

6

Bookmark Share ...



Clément Canonne @ccanonne.github.io · 4mo

In that case, this is very worrying -- there is no way the arXiv can do this at scale. Our entire *peer-review system* struggles to do that at scale.

(So I hope the policy is restricted to hallucinated references, which is something quite easy to detect)

2

2

11

Bookmark Share ...



Iris van Rooij @irisvanrooij.bsky.social · 4mo

I was responding to Jon's point about evidence.

Expert cognition can deliver inconvertible evidence of slop.

What I mean is, trying to automate slop detection is like trying to automate cognition. Not possible. That is not a failing of cognition or peer review, but a limitation of automation.

1

1

15

Bookmark Share ...



Iris van Rooij @irisvanrooij.bsky.social · 4mo

I see some smoking guns can be detected more easily 🤔 [bsky.app/profile/tdie...](#)

Blocked



[Add comment here later.]



1



Ted Dunning @bsky.tdunning.com · 4mo

How are you handling the cases of fake co-authors? How can I avoid being listed as a co-author on a paper that I had nothing to do with?

I haven't been hit by that yet, but I've heard from several famous people that they have been listed as co-authors on papers they had never heard of.



1



2



Thomas Dietterich @tdietterich.bsky.social · 4mo

If you discover that this has happened, please file a Code of Conduct report with ArXiv. The paper will be withdrawn and the author penalized.



1



3



Ted Dunning @bsky.tdunning.com · 4mo

Sure. Absolutely.

But how do you handle the penalties for the falsified authors?

In the case of authorship falsification, the submitter will know and should be subject to penalties, but the falsified authors likely had absolutely no knowledge of the paper being submitted. Will they be penalized?



1



1



Thomas Dietterich @tdietterich.bsky.social · 4mo

A weakness of our current system is that we do not have any identifying information about coauthors, beyond their names. They need not have ArXiv accounts, and they do not need to confirm co-authorship. This makes it difficult, but not impossible, to penalize them in practice



1



1



2



Thomas Dietterich @tdietterich.bsky.social · 4mo

But they should be penalized. They can then appeal and we can ask the submitter to produce evidence of consent to be coauthor. In the absence of such consent, the involuntary coauthor should not be punished.



1



1



4



Daniel Kane @daneelssoul.bsky.social · 4mo

Why would you expect the submitter provide such evidence even if it existed?



2



Bluesky

Sign in

(proving a negative is hard), so it is up to A to prove the presence of consent.

2

1

3



Thomas Dietterich @tdietterich.bsky.social · 4mo

I believe ArXiv normally withdraws the paper until the dispute is resolved.



2



Shoo away! What you're doing IS fascism @sahrataylor.bsky.social · 4mo

It should be a logarithmic progression

first ban is 1 year

Second year is 10 years

Third ban 100 years.

Yeah I know, but with AI there's no expiry date on identity.

And the names of the offending authors should be published!



2



howidiotami.bsky.social @howidiotami.bsky.social · 4mo

"...one year ban...." 🤡 why not 5 years? 10 years? bc a liar is a liar.. & will remain a liar.. misdirecting the efforts of thousands of others...



2



q-bert.bsky.social @q-bert.bsky.social · 4mo

I think it should be permanent. I think this garbage should be career-ending. This is some Marc Hauser nonsense -- lazier, apparently more tempting to many, but every bit as unethical.



1



David Rousseau @dhpmrou.bsky.social · 4mo

We'll soon see open source tools flourish to automatically check citations, and a few weeks later, such tools will be integrated in AI paper mills, and voilà, we'll still have more and more AI slop papers, but with no easy way to spot them.



Bruno @just--bruno.bsky.social · 4mo

Good call. It's wild how many papers have "hallucinated" citations now. Quality control is definitely needed.



Bluesky

Sign in

Something we can help with? We're picking up a lot of alerts in arXiv. Not all of them genAI related.



Nafnlaus 🇧🇪 🇪🇺 🇺🇦 @nafnlaus.bsky.social · 4mo

The ACL policy sounds reasonable. Yours does not. A 1 year ban with severe return requirements is the sort of penalty you'd get for like deliberately falsifying data or plagiarizing another paper. That's a penalty for severe academic misconduct, not carelessness.

1



Nafnlaus 🇧🇪 🇪🇺 🇺🇦 @nafnlaus.bsky.social · 4mo

And you apply the penalty not just for hallucinated references, but for things that are even merely "misleading" or "biased".

Since when is it a normal penalty for an author who lets a "Insert figure here" slide through to get such harsh penalties?

2



Saxxie.dev @saxxie.dev · 4mo

It's extremely easy to not fall afoul of this policy - just read the entire paper. if that's too difficult - pretty much any modern LLM can reliably double check whether you missed anything. If you can't figure out how to get an LLM to proofread you probably shouldn't be using them to write.

1



6



Nafnlaus 🇧🇪 🇪🇺 🇺🇦 @nafnlaus.bsky.social · 4mo

Are you going to apply that same standard to human errors?

"It's extremely easy not to fall afoul of having a "Insert figure here" slip through - just read through the entire paper"

With human errors, we treat them as minor. Just upload a v2. With LLM errors, they're treating it as career-ending.

3



Saxxie.dev @saxxie.dev · 4mo

with a human error, it means you only read it once (while writing it). with a machine error, it means you read it zero times. This is a huge difference! (infinty percent, in the one direction)



1



Bluesky

Sign in

