



A Retrospective on the ICLR 2026 Review Process

[ICLR 2026 PROGRAM CHAIRS](#) / [ICLR 2026](#)

The selection of papers for ICLR 2026 has fully concluded. We extend our congratulations to the authors whose work will appear at the conference. Creating ICLR's technical program requires immense effort from the authors, reviewers, and area chairs, and we thank you for your contributions and service. For researchers whose work was rejected, we hope that the review process will help you improve your submission for future publication.

ICLR 2026 received 19,525 valid, format-compliant submissions. 779 submissions were desk rejected for procedural/content violations and 5,042 were withdrawn, so that ultimately 13,763 submissions received an accept or reject decision based on 76,139 reviews provided by 18,054 reviewers. These numbers reflect a significant increase in submissions from past years. Of those decisions, 5,355 papers were accepted and 8,408 were rejected, representing an acceptance rate of 27.4%.

The ICLR 2026 program also faced unique challenges due to the rise of LLMs in submissions/reviews and the security incident. Below, we summarize these challenges and our responses.

Large Language Models

The capabilities of AI systems based on large language models (LLMs) have grown over time and the usage of LLMs in the publication and peer review process has grown in tandem. We proactively designed a set of [policies](#) on the usage of LLMs based on ICLR's [Code of Ethics](#). In short, these policies stipulated

that 1) any usage of LLMs needed to be disclosed and 2) individuals were ultimately responsible for their contributions. Violations of these policies were treated as a code of ethics violation. While we considered any violation brought to our attention, our enforcement of these policies was largely achieved by flagging reviews that were likely generated by an LLM and by finding references in submissions that referred to nonexistent publications.

Dealing with LLM-Generated Reviews

There are many ways a reviewer might have used an LLM in their review, such as refining their writing, summarizing a submission, or even writing a review wholesale. While systems exist for detecting LLM-generated content, their accuracy is far from perfect, and we did not want to penalize cases where an LLM was used to assist with writing but the judgements made in the review were reasonable and valid. On the other hand, we needed to proactively address the possibility that many were using LLMs to offload their reviewing responsibilities, ultimately producing sub-par reviews. The large scale of reviews (over 75,000) made it impossible to rely on human judgement alone to consistently flag and discover problematic reviews. We consequently ran two LLM content detectors on all submitted reviews and emailed area chairs to notify them of which reviews on their submissions were flagged as being entirely LLM generated by both detectors. Area chairs were then asked to consider this analysis as an aspect of review quality when flagging low-quality reviews. Ultimately, this approach is in line with standard practice for flagging possible low-quality reviews (e.g. those that are excessively short) to area chairs. As in past years, reviewers who were repeatedly flagged as writing low-quality reviews will be removed from the reviewing pool going forward.

Taking Actions on Hallucinated References

As in reviews, there are many ways an LLM could be used in the preparation of an ICLR submission, some of which might improve its quality and others that might be harmful. Similar to our approach to reviews, we ran automated LLM content detection on all submissions and flagged submissions with a high proportion of LLM-generated content to ACs. Additionally, we found one widespread issue stemming from LLM use that was relatively straightforward to

detect: hallucinated references. Specifically, we found that many submissions included one or more references that referred to documents that didn't exist and/or contained egregiously incorrect bibliographic information. To automatically detect these cases, we made use of a system that automatically extracted references from a given submission and checked them against multiple standard bibliographic databases and a standard web search. This system had a significant false positive rate (for example, it would flag a reference to a paper with a non-English title that a submission's authors had translated to English), so we relied on area chairs to perform a first round of human review of flagged references. Then, we (the program chairs) checked all flagged references manually ourselves. Ultimately, each paper with flagged references was reviewed by at least three humans. Finally, all papers with confirmed hallucinated references were desk rejected, together with an appeal channel to take care of any erroneously flagged papers. This process partially explains the relatively high desk rejection rate at this year's ICLR compared to past years.

OpenReview Security Incident

In the middle of the discussion period, we had an unprecedented disruption to the review process: A malicious user [exploited the openreview API](#) to scrape, and ultimately release, the identities of the authors, reviewers, and area chairs for a large subset of ICLR 2026 submissions. Immediately afterward, we started receiving reports of collusion attempts, threats, and harassment, all leading to reviewers feeling pressured to change their scores. A summary of our response to this incident can be found [here](#). Our primary goal was to preserve the integrity of the reviews process. Because the underlying vulnerability was introduced at the beginning of the rebuttal phase three weeks prior to the scrape, we reset all review scores to the pre-rebuttal state, froze the discussion and reassigned area chairs to each submission. The new area chairs were then tasked with inferring the expected outcome had discussion proceeded as normal, exercising their judgement in interpreting the discussion that had already happened. To give area chairs time to complete this increased workload, we substantially extended the meta-review period. After decisions came in, we found the overall trend of acceptance based on reviewer opinions was similar prior years, suggesting that this approach was effective. We

acknowledge the additional stress this caused; we did not take the decision lightly but chose this course of action after careful deliberation among many unideal options.

Apart from modifying the meta-review process, we investigated all reported cases of collusion, threats, and harassment. While many cases originated from individuals from outside the ICLR community, we ultimately banned and desk rejected the submissions of any offending members of the ICLR community. The malicious user who exploited the openreview API was also banned by OpenReview.

Reflections

Despite the large increase in scale and the unexpected twists and turns, we ultimately were able to produce a program that we consider to reflect ICLR's established standards. Our efforts were supported by the large community of program committee members (especially area chairs!) who undertook an unusually large workload this year. We hope that the remediation efforts can help inform policies and approaches used at future conferences.

[!\[\]\(e78f798d4ea5c530c9db49e7d26e6b95_img.jpg\) Workshops at ICLR 2026](#)

[Announcing the ICLR 2026 keynotes !\[\]\(23d9fc146e83b5c3013cfa32c784f8d5_img.jpg\)](#)

Related Posts

ICLR 2026

**Announcing the ICLR
2026 Outstanding
Papers**

ICLR 2026

**Announcing the Test
of Time Awards from
ICLR 2016**

ICLR 2026

**Announcing the ICLR
2026 keynotes**

 Search

Recent Posts

[Making Google's Paper Assistant Tool \(PAT\) available to ICLR submitters](#)

SEPTEMBER 10, 2026

[Submission policies for ICLR 2027](#)

SEPTEMBER 2, 2026

[Announcing the ICLR 2026 Outstanding Papers](#)

APRIL 23, 2026

[Announcing the Test of Time Awards from ICLR 2016](#)

APRIL 22, 2026

[Announcing the ICLR 2026 keynotes](#)

APRIL 17, 2026

[A Retrospective on the ICLR 2026 Review Process](#)

MARCH 31, 2026

[Workshops at ICLR 2026](#)

JANUARY 13, 2026

[ICLR's Commitment to OpenReview](#)

DECEMBER 17, 2025

[ICLR 2026 Response to Security Incident](#)

DECEMBER 3, 2025

[ICLR 2026 Response to LLM-Generated Papers and Reviews](#)

NOVEMBER 19, 2025

[Policies on Large Language Model Usage at ICLR 2026](#)

AUGUST 26, 2025

[ICLR 2025 Mentoring Chats](#)

APRIL 23, 2025

Recent Comments
